nODD Gen3 — derived pT comparison

PROTOTYPE · Fatras · constant 3 T · all −4 ≤ η ≤ 4 · 20,000 generated muons per momentum

Core ±5; density normalized to all finite entries, including those outside this view.

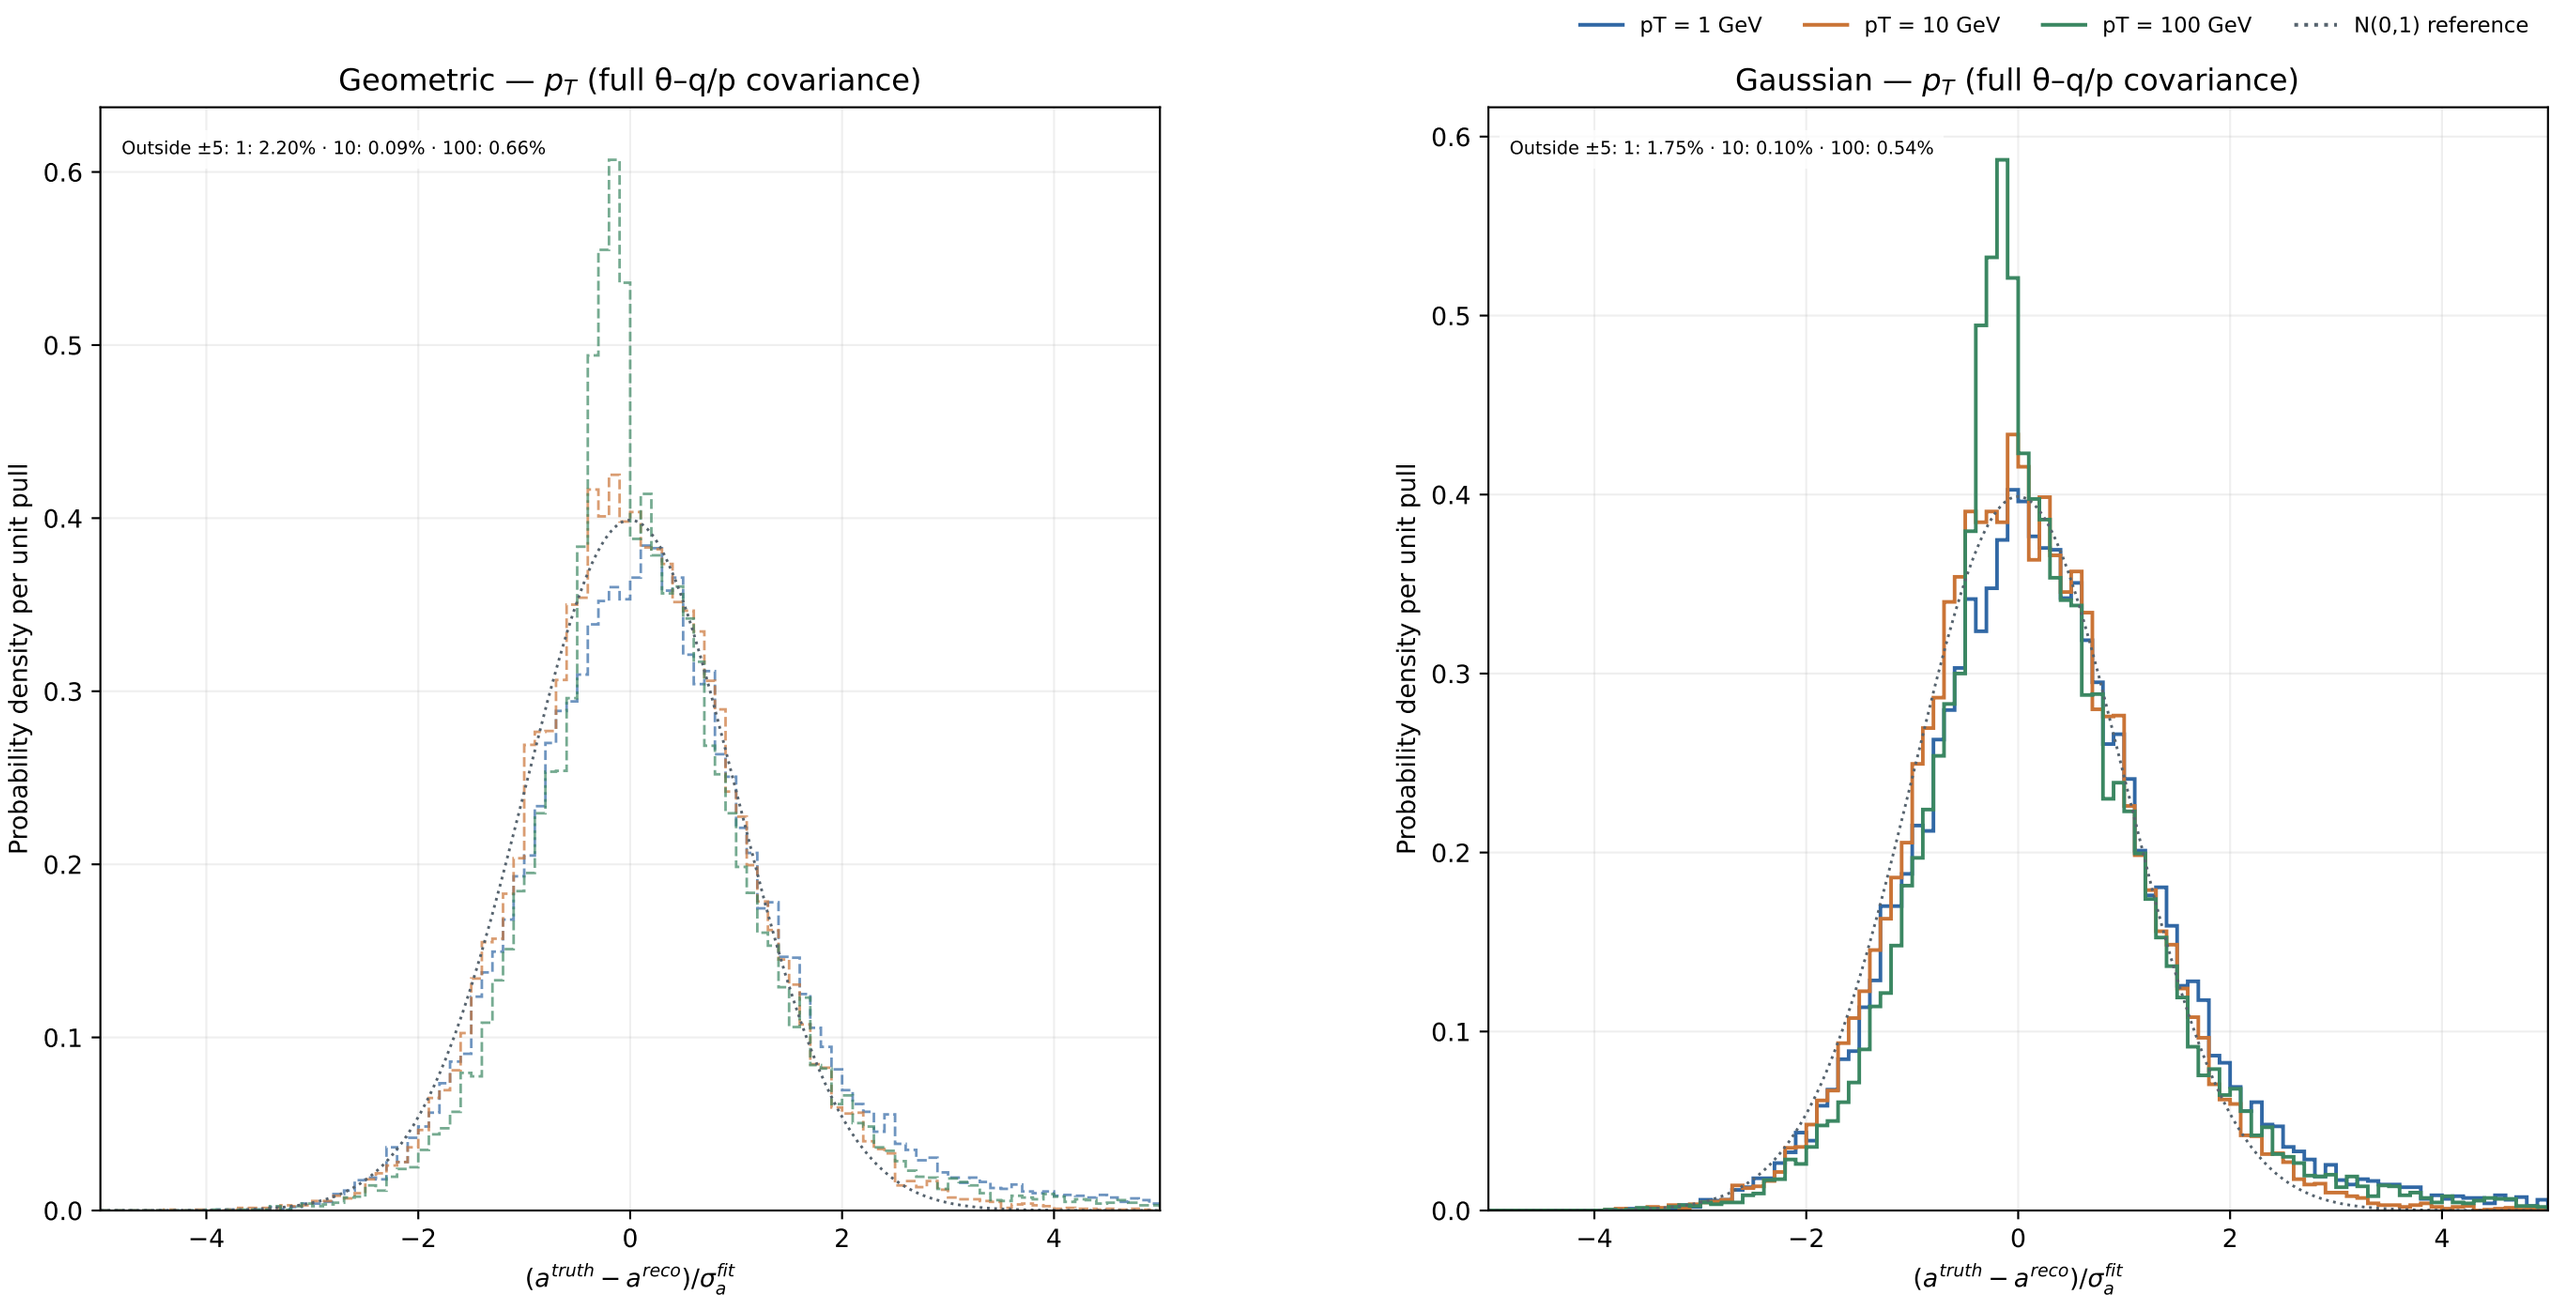

All training cluster sizes pooled by volume/coordinate; no fitted-pull tuning. Native out-of-bounds Gaussian draws discarded without retry. Time is unmeasured.